

Daniel Thomas Jesudoss

Dallas - Fort Worth, TX | idanycb@gmail.com | linkedin.com/in/danycb | github.com/idanycb | danycb.com

Software engineer with an MS in Computer Science from UT Arlington (May 2026). Built and run two live projects solo: FinDoc, a RAG app that answers questions about SEC filings with citations, and the K3s cluster on Oracle Cloud that hosts it. Most comfortable in Java, Python and TypeScript.

Education

Master of Science, Computer Science | GPA: 3.9

May 2026

The University of Texas at Arlington

Arlington, TX

- Relevant Coursework: Design and Analysis of Algorithms, Software Engineering, Intelligent Systems

Bachelor of Technology, Information Technology

Apr 2024

Loyola-ICAM College of Engineering and Technology

Chennai, India

Technical Skills

Languages: Java, Python, TypeScript, JavaScript, SQL

Databases: PostgreSQL, pgvector (HNSW), Flyway migrations

Testing: JUnit, Mockito, Testcontainers, pytest

Frameworks: Spring Boot, React, Next.js, Flask, FastAPI, Node.js

Cloud & DevOps: AWS (S3, SQS, EC2, Lambda, Route 53), Oracle Cloud, Docker, Kubernetes, Flux CD, GitHub Actions

AI & Architecture: RAG, LangChain4j, Gemini, Groq, prompt engineering, hexagonal architecture, Claude Code, Cursor

Experience

IBM | Frontend Development Intern | Remote

Sep 2022 – Nov 2022

- Built an internal dashboard in React and TypeScript that pulled live team metrics from IBM REST APIs into interactive charts and tables.
- Wrote TypeScript scripts to clean and prepare the data behind the dashboard's charts, working mostly on my own.
- Took the project from scoping to a live demo for management in IBM's three-month SkillsBuild program.

Projects

FinDoc Analyzer | Java 21, Spring Boot, pgvector, LangChain4j, Gemini, AWS S3/SQS, Next.js

Jan 2026

- Built a self-hosted RAG app that imports 10-K and 10-Q filings from SEC EDGAR, tracks which sections each amendment replaces, and answers questions with numbered citations back to the filing.
- Designed retrieval as up to three passes (vector search, LLM query rewrite, HyDE fallback) over 384-dim MiniLM embeddings in pgvector. Docling parses uploads, and S3 events on SQS drive async ingestion.
- Kept each team's documents isolated with team-scoped queries and three-role RBAC over JWT auth. Structured the Spring Boot backend as ports-and-adapters, with JUnit and Testcontainers tests covering isolation and auth.
- Built an eval harness on a hand-checked gold set from Tesla's 2025 filings: evidence recall@6 0.90, MRR 0.92, amendment version accuracy 1.00.

Portfolio GitOps | K3s, Flux CD, Infisical, Traefik, OCI ARM

Feb 2026

- Run a single-node K3s cluster on Oracle Cloud's free ARM tier that hosts danycb.com and FinDoc. Everything after VM creation comes from Git, with no SSH.
- Broke the bootstrap loop (Flux needs a key stored in Infisical, which needs a running cluster) with a two-phase cloud-init script, so the setup rebuilds on a fresh VM.
- Infisical serves secrets through the cluster's Kubernetes identity, so the repo has no static credentials. Flux syncs Git every minute, reverts drift, and rolls out new GHCR images on its own.

Instagram Sentiment Analyzer | Final-year team project | Python, Flask, Next.js, TypeScript, scikit-learn

May 2024

- Built the Flask REST API (5 endpoints) for Instagram login, post and comment fetching, account stats and sentiment breakdowns, using Instagrapi.
- Plugged a pre-trained emotion classifier into the comment pipeline and charted the results in a Next.js + TypeScript dashboard with Material-UI.